

COT 6405
ANALYSIS OF ALGORITHMS

Dynamic Programming

Computer & Electrical Engineering and Computer Science Department
Florida Atlantic University

Outline

- DP – introduction
- Weighted Interval Scheduling (KT – chapter 6.1)
- Principles of DP (KT – chapter 6.2)
- Change-making problem (BB – chapter 8.2)
- 0-1 knapsack problem (BB – chapter 8.4)
- Sequence alignment problem (KT – chapter 6.6)

KT book - *Algorithm Design* by J. Kleinberg and Eva Tardos

BB book - *Fundamentals of Algorithms* by Gilles Brassard and Paul Bratley

Dynamic Programming (DP)

- is a technique, not a specific algorithm (like divide-and-conquer)
- applied to optimization problems (maximization or minimization)
- applicable when subproblems are not independent, that is subproblems share subsubproblems. Then DP solves each subproblem only once, stores the result in a table, and reuses it later

Weighted Interval Scheduling

Problem definition:


We are given **1 resource** and a number of **n requests** labeled $1, 2, \dots, n$. Each request i has a starting time s_i , a finishing time f_i , and a *value* (or *weight*) $v_i > 0$.

Find a *compatible* subset S of requests (intervals) of *maximum total value* $\sum_{i \in S} v_i$.

- two requests i and j are compatible if the requests do not overlap
 - request i is earlier than j , $f_i \leq s_j$
 - request j is earlier than i , $f_j \leq s_i$
- a subset of requests is compatible if all pairs i, j ($i \neq j$) are compatible


Weighted Interval Scheduling

Example



Weighted Interval Scheduling


- can be solved using *greedy* when all requests have the same value (equal to 1)
 - greedy choice: choose the compatible request with the earliest finishing time
- when the requests have different values, greedy does not work



- for different request values, DP produces an optimal solution

A recursive algorithm

- suppose that the requests are sorted in nondecreasing finishing time $f_1 \leq f_2 \leq \dots \leq f_n$. We say that request i *comes before* j if $i < j$.
- $p(j)$ – the largest index $i < j$ such that requests i and j are disjoint
 - rightmost interval i that ends before j begins



Weighted Interval Scheduling

- optimal solution O : either n belongs to O or it doesn't
 - if $n \in O$, then in addition O contains an *optimal* solution to the problem with requests $\{1, \dots, p(n)\}$
 - if $n \notin O$, then O is an *optimal* solution to the problem with requests $\{1, 2, \dots, n-1\}$
- let O_j – optimal solution for the problem w/ requests $\{1, 2, \dots, j\}$
 - $OPT(j)$ – value of this solution
- objective: find O_n and $OPT(n)$

Writing a recursion

$$\text{OPT}(j) = \max (v_j + \text{OPT}(p(j)), \text{OPT}(j-1))$$

- request j belongs to an optimal solution on the set $\{1, 2, \dots, j\}$ iff $v_j + \text{OPT}(p(j)) \geq \text{OPT}(j-1)$
- characteristic of DP: write a recurrence equation that expresses the optimal solution (or its value) in terms of optimal solutions to smaller subproblems

Recursive algorithm

Compute-Opt(j)

if $j == 0$ then

 return 0

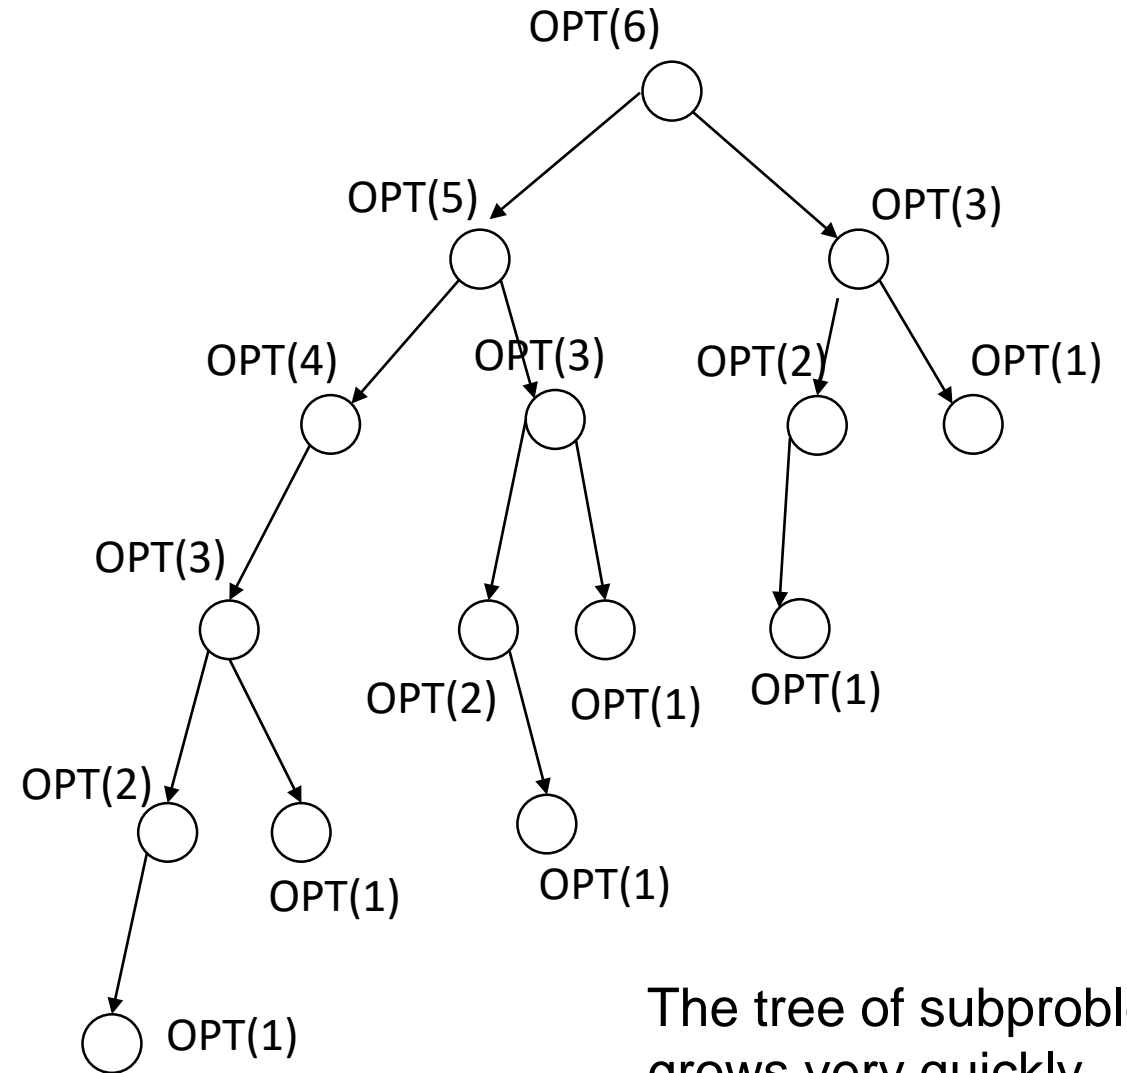

else

 return $\max(v_j + \text{Compute-Opt}(p(j)), \text{Compute-Opt}(j-1))$

endif

- drawback: RT is exponential in n

Example



The tree of subproblems grows very quickly.

Memoizing the Recursion

- Key observations:
 - Compute-Opt(n) is only solving $n+1$ different subproblems:
Compute-Opt(0), Compute-Opt(1), ..., Compute-Opt(n)
 - exponential time is due to the redundancy in the number of times the same call is made
- **Memoization technique:** *solve each subproblem only once and store its value in a table. All future calls use the precomputed value.*

Memoizing the Recursion

- array $M[0..n]$
 - $M[j]$ initially empty, then store $\text{Compute-Opt}(j)$ value
- to determine $\text{OPT}(n)$, call $\text{M-Compute-Opt}(n)$

M-Compute-Opt(j)

if $j == 0$ then

 return 0

elseif $M[j]$ is not empty then

 return $M[j]$

else

$M[j] = \max(v_j + \text{M-Compute-Opt}(p(j)), \text{M-Compute-Opt}(j-1))$

 return $M[j]$

endif

Analyzing the Memoized Version

- Assuming that the input intervals are sorted by their finish time and that the $p()$ values are already computed, the RT to compute $M\text{-Compute-Opt}(n)$ is $O(n)$.
 - Proof: since M has $n+1$ entries, there are at most $O(n)$ calls to $M\text{-Compute-Opt}$, and hence the RT is $O(n)$

A second solution: iterative approach

Iterative-Compute-Opt

$M[0] = 0$

for $j = 1, 2, \dots, n$

$M[j] = \max(v_j + M[p(j)], M[j-1])$

endfor

$RT = O(n)$

- Example

Computing a solution in addition to its value

- Find the subset of requests of maximum value

Find-Solution(j)

if $j == 0$ **then**

 return \emptyset

else

if $v_j + M[p(j)] \geq M[j-1]$ **then**

 output $\text{Find-Solution}(p(j)) \cup \{j\}$

else

 output the result of $\text{Find-Solution}(j-1)$

endif

endif

- initial call: $\text{Find-Solution}(n)$
- $RT = O(n)$, calls recursively on smaller instances and takes constant time per call

Principles of Dynamic Programming

- write solution to a problem recursively, based on solutions to subproblems
- memoization or iteration over subproblems?
 - iteratively building up subproblems is simpler
- properties of DP:
 - there is only a polynomial number of subproblems
 - the solution to the original problem can be easily computed from the solutions to the subproblems
 - there is a natural ordering of subproblems from smallest to largest

Change-making problem

- Unfortunately, even though greedy algorithm is very efficient, it works only in a limited number of instances
- Dynamic programming works for all systems of coins
- Suppose that the currency has available coins of n different denominations
 - a coin of denomination i , $1 \leq i \leq n$ has value $d_i > 0$ units
- Suppose that we have an unlimited supply of coins of each denomination
- Goal: give the customer coins worth N units, using as few coins as possible

DP approach

- table $c[1..n, 0..N]$ – one row for each denomination and one column for each amount from 0 units to N units
- $c[i,j]$ – minimum number of coins required to pay an amount j , with $0 \leq j \leq N$, using only coins of denominations 1 to i , $1 \leq i \leq n$
- the solution to the original problem is given by $c[n, N]$

Filling up the table

- $c[i,0] = 0$ for all i
- then the table can be filled row by row, left to right, or column by column, top to bottom
- to compute $c[i,j]$, two choices:
 - do not use any coins of denomination i , then $c[i,j] = c[i-1,j]$
 - use at least one coin of denomination i , then $c[i,j] = 1 + c[i,j - d_i]$
- therefore:
$$c[i,j] = \min(c[i-1,j], 1+c[i,j-d_i])$$
- if $i = 1$ and $j < d_1$, then set $c[i,j] = \infty$

DP algorithm

- Algorithm “coins” computes the minimum number of coins needed to make change for N units.
- The array d[1..n] specifies the coinage. In the example, there are coins for 1, 4, and 6 units.

Algorithm coins(N)

array d[1..n] = [1, 4, 6]

array c[1..n, 0..N]

for i = 1 to n

 c[i, 0] = 0

for i = 1 to n

for j = 1 to n

if i == 1 and j < d[i] then c[i, j] = ∞

elseif i == 1 then c[i, j] = 1 + c[1, j - d[1]]

elseif j < d[i] then c[i, j] = c[i-1, j]

else c[i, j] = min(c[i-1, j], 1 + c[i, j - d[i]])

return c[n, N]

RT = $\Theta(nN)$

DP algorithm comments

- if an unlimited supply of coins of value 1 is available, then we can always find a solution to our problem
- if no coin with value 1, then there may be values of N for which no solution is possible
 - algorithm returns ∞

Finding how many coins of each denominations are used

PrintCoins(c,i,j)

```
if c[i,j] =  $\infty$ 
    then return no change possible
if j == 0 then return
if c[i,j] == c[i-1,j]
    then PrintCoins(c,i-1,j)
elseif c[i,j] == 1 + c[i,j-di]
    then PrintCoins(c,i,j-di)
    print di
```

- Initial call: PrintCoins(c,n,N)
- RT = $O(N + n)$
- Example

0-1 knapsack problem

- Given:
 - n objects and a knapsack
 - for $i = 1, \dots, n$, object i has a positive weight w_i and a positive value v_i
 - objects may not be broken into smaller pieces, either take the whole object or leave it behind
 - the knapsack can carry a weight $\leq W$
- Objective: fill the knapsack s.t. to maximize the value of the included objects, while respecting the capacity constraints

DP approach

- table $V[1..n, 0..W]$ – one row for each available object and column for each weight from 0 to W
- $V[i, j]$ – maximum value of the objects we can transport if the weight limit is j , $0 \leq j \leq W$, and if we only include objects numbered from 1 to i , $1 \leq i \leq n$
- The solution to the original problem is given by $V[n, W]$

Filling up the table

- $V[i,0] = 0$, for all i
- then the table can be filled row by row, left to right, or column by column, top to bottom
- to compute $V[i,j]$, two choices:
 - not adding object i to the knapsack, $V[i,j] = V[i-1, j]$
 - adding object i to the knapsack, $V[i,j] = v_i + V[i-1, j-w_i]$
- therefore:
$$V[i,j] = \max(V[i-1,j], v_i + V[i-1,j-w_i])$$
- for the out-of-bounds-entries, we define:
$$V[0,j] = 0 \text{ when } j \geq 0$$
$$V[i,j] = -\infty \text{ for all } i \text{ when } j < 0$$

Sequence Alignment – first example

- dictionaries on the Web and spell checkers
 - if you type “*ocurrence*”, it may ask: “Perhaps you mean *occurrence*?”
 - the dictionary will search its entries for the word most “similar” to the one typed
- how should we define *similarity* between two words or strings?

Modeling “similarity”

- model similarity between two strings by the number of *gaps* and *mismatches* that occur when lining up the two sequences

o	-	c	u	r	r	a	n	c	e
o	c	c	u	r	r	e	n	c	e

- one gap
- one mismatch

o	-	c	u	r	r	-	a	n	c	e
o	c	c	u	r	r	e	-	n	c	e

- three gaps
- no mismatch

- which one is better?

Computational biology – second example

- An organism's *genome*:
 - divided into linear DNA molecules, chromosomes, which serve as an one-dimensional chemical storage device
 - string over the alphabet {A,C,G,T} that determines the properties of the organism
 - adenine, cytosine, guanine, thymine
 - the string encodes instructions for building protein molecules
 - using a chemical mechanism to read portions of the chromosomes, a cell can construct proteins that control its metabolism

Computational biology – second example, cont.

- let X and Y be two strains of bacteria, closely related evolutionary
- assume that a certain substring in the DNA of X codes for a certain toxin
- if this substring is found in Y , we may hypothesize that this portion codes for a similar kind of toxin

Sequence Alignment Problem

- let $X = x_1x_2..x_m$ and $Y = y_1y_2..y_n$
- $\{1,2,...,m\}$ and $\{1,2,...,n\}$ represent different positions in strings X and Y
- *matching*: set of ordered pairs such that each item occurs in at most one pair
- a matching M of these two sets is an *alignment* if there are no *crossing* pairs: if $(i,j), (i',j') \in M$ and $i < i'$ then $j < j'$
 - an alignment provides a way to line up the two strings

stop- -tops

- alignment: $\{(2,1),(3,2),(4,3)\}$

Sequence Alignment Problem

Let M be an alignment of X and Y

- *gap penalty*: each gap incurs a cost $\delta > 0$
- *mismatch cost*: for each pair of letters p, q in the alphabet, there is a mismatch cost α_{pq} for lining up p with q
 - usually $\alpha_{pp} = 0$
- the cost of M is the sum of gap and mismatch costs

Objective: find an **optimal alignment** of X and Y , that means an alignment of minimum cost.

- values δ and $\{\alpha_{pq}\}$ are given parameters
- the lower the cost, the more similar X and Y are
- going back to the first example *ocurance* and *occurrence*, the first alignment is better if and only if $\delta + \alpha_{ae} < 3\delta$

Designing the DP algorithm

- in the optimal alignment M :
 - either $(m,n) \in M$ or $(m,n) \notin M$
- in any alignment M :
 - if $(m,n) \notin M$, then either the m^{th} position of X or the n^{th} position of Y are not matched in M
 - proof: assume by contradiction $(i,n), (m,j) \in M$ for some i,j .
Then $i < m$ and $j < n \Rightarrow$ crossing pairs \Rightarrow contradiction

Property: in any alignment M , one of the following is true:

- (i) $(m,n) \in M$, or
- (ii) the m^{th} position of X is not matched, or
- (iii) the n^{th} position of Y is not matched

Designing the DP algorithm

- let $\text{OPT}(i,j)$ – minimum cost of an alignment between $x_1x_2\dots x_i$ and $y_1y_2\dots y_j$
- recursively define $\text{OPT}(m,n)$:

case (i):

- pay $\alpha_{x_my_n}$, then optimally align $x_1x_2\dots x_{m-1}$ and $y_1y_2\dots y_{n-1}$
 $\text{OPT}(m,n) = \alpha_{x_my_n} + \text{OPT}(m-1,n-1)$

case (ii):

- pay gap cost δ , then optimally align $x_1x_2\dots x_{m-1}$ and $y_1y_2\dots y_n$
 $\text{OPT}(m,n) = \delta + \text{OPT}(m-1,n)$

case (iii):

- pay gap cost δ , then optimally align $x_1x_2\dots x_m$ and $y_1y_2\dots y_{n-1}$
 $\text{OPT}(m,n) = \delta + \text{OPT}(m,n-1)$

Designing the DP algorithm

Property: The minimum alignment costs satisfy the following recurrence, for $i \geq 1, j \geq 1$:

$$\text{OPT}(i,j) = \min\{\alpha_{x_i y_j} + \text{OPT}(i-1,j-1), \delta + \text{OPT}(i-1,j), \delta + \text{OPT}(i,j-1)\}$$

Moreover, (i,j) is in an optimal alignment iff the minimum is achieved by the first of these values.

Alignment (X,Y)

array $A[0..m,0..n]$

initialize array $A[i,0] = i\delta$ for each i

initialize array $A[0,j] = j\delta$ for each j

for $j = 1$ to n

for $i = 1$ to m

$$A[i,j] = \min \{ \alpha_{x_i y_j} + A[i-1,j-1], \delta + A[i-1,j], \delta + A[i,j-1] \}$$

return $A[m,n]$

$$\text{RT} = \Theta(mn)$$