

COT 6405
ANALYSIS OF ALGORITHMS

Greedy Algorithms

Computer & Electrical Engineering and Computer Science Department
Florida Atlantic University

Outline

- Greedy algorithms
- Problems solved using greedy
 - Scheduling all intervals (KT – chapter 4.1)
 - Scheduling to minimize lateness (KT – chapter 4.2)
 - Change making problem (CLRS-problem 16-1 page 446)
 - The knapsack problem (CLRS page 425-427)

KT book - *Algorithm Design* by J. Kleinberg and Eva Tardos

CLRS book - *Introduction to Algorithms*, 3rd edition, by T. H. Cormen, C. E. Leiserson, R. L. Rivest, and C. Stein



Greedy Algorithms

- Used for optimization problems
- When we have to make a choice, make the choice that looks best at the moment
- A greedy algorithm runs over a number of steps: at each step we make a greedy choice and we are left with one subproblem to solve
- A greedy algorithm does not necessarily produce an optimal solution. We need to prove it.

Problem: Scheduling All Intervals

- Interval Partitioning Problem: we have many identical resources available and we want to schedule all the requests using as few resources as possible
- Example:
 - Each request is a lecture to be scheduled in a classroom for a particular interval of time
 - Objective: satisfy all the requests using as few classrooms as possible
 - Constraints: any two classes that overlap in time must be scheduled in different classrooms

Example



- 10 intervals (*a* through *j*)
- all intervals can be scheduled using 3 resources: each row represents a set of intervals that can be scheduled on a single resource

Interval Partitioning Problem

- Define *depth* of a set of intervals as the maximum number of intervals that pass over a single point on the time-line
- Property: the number of resources needed is at least the depth of the set of intervals
- Design a greedy algorithm that schedules all intervals using a number of resources equal to the depth
 - Optimality of the algorithm results from the property

Greedy algorithm

- Let d – depth of the set of intervals

Algorithm

sort the intervals by their start times, breaking ties arbitrarily

let I_1, I_2, \dots, I_n denote the intervals in the sorted order

compute depth d

for $j = 1, 2, 3, \dots, n$

for each interval I_i that preceded I_j in the sorted order and overlaps it
 exclude the label of I_i from consideration for I_j

if there is any label from $\{1, 2, \dots, d\}$ that has not been excluded
 assign a nonexcluded label to I_j

else

 leave I_j unlabeled

Analyzing the algorithm

- Property: using the greedy algorithm, every interval will be assigned a label, and no two overlapping intervals will receive the same label

Proof:

consider interval I_j , assume there are t intervals earlier in the sorted order that overlap it
 $t+1$ overlapping intervals $\Rightarrow t+1 \leq d \Rightarrow t \leq d - 1 \Rightarrow$ there is at least one label available

No two overlapping intervals receive the same label: the second interval in the list will be assigned a different label

- Property: the proposed greedy algorithm schedules each interval on a resource, using a number of resources equal to the depth of the set of intervals. This is the optimal number of resources needed.

$$RT = O(n^2)$$

Outline

- Greedy algorithms
- Problems solved using greedy
 - Scheduling all intervals (KT – chapter 4.1)
 - **Scheduling to minimize lateness (KT – chapter 4.2)**
 - Change making problem (CLRS-problem 16-1 page 446)
 - The knapsack problem (CLRS page 425-427)

Problem: scheduling to minimize lateness

Problem description:

- a single resource
- a set of n requests, where each request i has a deadline d_i and requires a contiguous time interval of length t_i
- different requests must be assigned nonoverlapping intervals
- each request will be satisfied, request i will be scheduled $[s(i), f(i)]$, where $f(i) = s(i) + t_i$
- a request i is *late* if it misses the deadline, $f(i) > d_i$
 - **lateness** defined as $\ell_i = f(i) - d_i$
 - if $\ell_i = 0$, then the request i is not late

Goal: schedule all requests (e.g. compute $s(i)$, $f(i)$ for all $i=1..n$), using nonoverlapping intervals, such that to minimize the maximum lateness,


$$L = \max_i \ell_i$$

Example

length = 1, deadline = 2

Job 1 

length = 2, deadline = 4


Job 2 

length = 3, deadline = 6

Job 3 

- First job has length $t_1 = 1$ and deadline $d_1 = 2$
- Second job has length $t_2 = 2$ and deadline $d_2 = 4$
- Third job has length $t_3 = 3$ and deadline $d_3 = 6$

Possible schedules:



Maximum lateness = 2



Maximum lateness = 0

- Optimal solution: maximum lateness is 0

What greedy choice to choose?

Several greedy choices are possible for requests (t_i, d_i) :

- Schedule jobs in the order of **increasing length t_i**
 - does not always lead to an optimal solution
 - example: two jobs $J1(t_1=1, d_1=100)$ and $J2(t_2=10, d_2=10)$
scheduling by increasing length: $J1 + J2 \Rightarrow L = 1$
optimal scheduling: $J2 + J1 \Rightarrow L = 0$
- Schedule jobs in the order of **increasing slack $d_i - t_i$**
 - does not always lead to an optimal solution
 - example: two jobs $J1(t_1=1, d_1=2)$ and $J2(t_2=10, d_2=10)$
scheduling by increasing slack: $J2 + J1 \Rightarrow L = 9$
optimal scheduling: $J1 + J2 \Rightarrow L = 1$
- Schedule jobs in the order of **increasing deadline d_i**
 - Always yields an optimal solution!

Greedy choice: earliest deadline first

- Greedy choice that always produce an optimal solution:
 - sort the jobs in **increasing order of their deadlines d_i**
 - schedule jobs in this order
- Assume that jobs are labeled in the order of their deadlines (rename them if necessarily)

$$d_1 \leq d_2 \leq \dots \leq d_n$$

- J1 starts at s and ends at $f(1) = s(1) + t_1$
- J2 starts at $f(1)$ and ends at $f(2) = s(2) + t_2$
- ... so on

Greedy choice: earliest deadline first

Algorithm:

- order the jobs in nondecreasing order of their deadlines

- assume for simplicity of notation that $d_1 \leq d_2 \leq \dots \leq d_n$

- initially $f = s$

- for each job in the sorted order

 - assign job i to the interval $s(i) = f$ and $f(i) = f + t_i$

 - $f = f + t_i$

- return the set of scheduled intervals $[s(i), f(i)]$ for $i=1 \dots n$

RT = $O(n \lg n)$

Correctness: using an exchange argument

- Observation:
 - our algorithm produces a schedule with no idle time
 - there is an optimal schedule with no idle time
- ***Exchange argument method***: start with an optimal solution and gradually modify it, preserving its optimality at each step, transforming it to the solution returned by the greedy algorithm
- Apply the “exchange argument” method to our problem:
 - let O be an optimal schedule
 - let A be the schedule returned by the greedy algorithm

Correctness

- a schedule A' has an inversion if for some $d_j < d_i$, the job i is scheduled before the job j
- the schedule A (greedy algorithm) has no inversion
- Property: ***all schedules with no inversions and no idle times have the same maximum lateness***
 - no inversions & no idle times \Rightarrow schedules can differ in the order in which jobs with identical deadlines are scheduled
 - all these schedules have the same maximum lateness
- Property: ***there is an optimal schedule that has no inversions and no idle time***

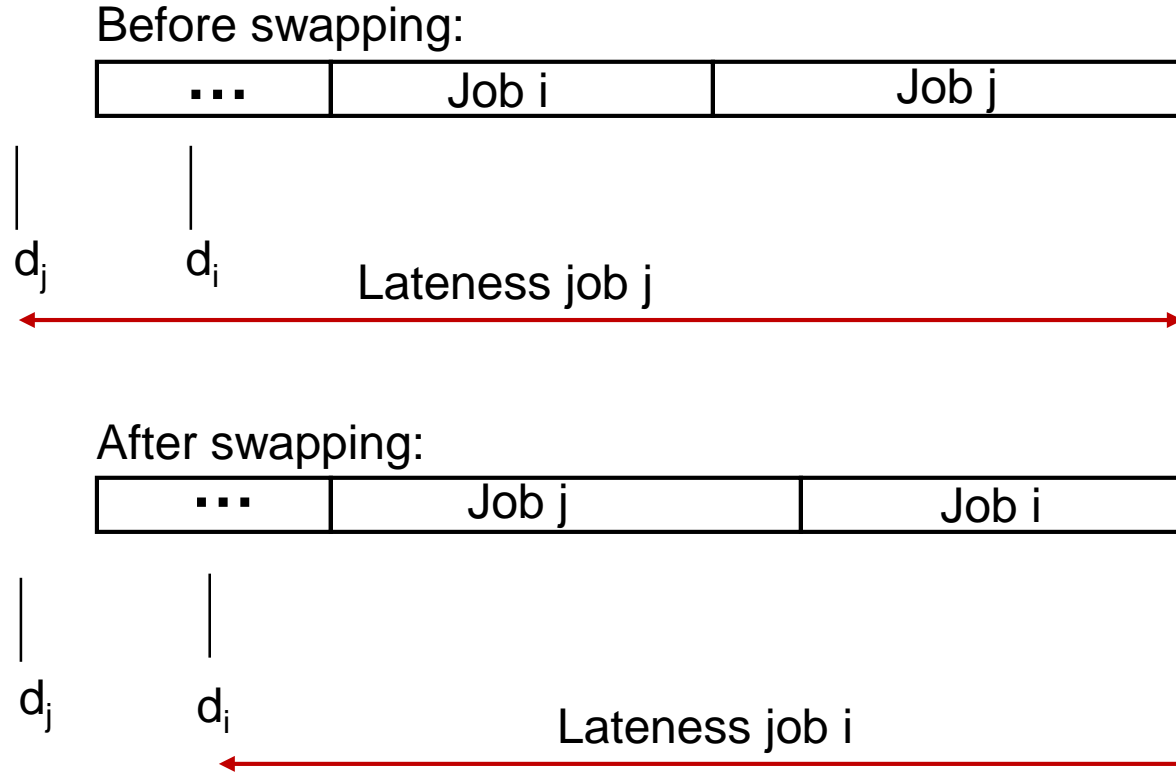
Correctness

Property: ***there is an optimal schedule that has no inversions and no idle time***

Proof:

- If O has an inversion, then there is a pair of jobs i and j such that j is scheduled immediately after i and has $d_j < d_i$
 - Examine the schedule starting from the beginning. At some point, the deadline decreases for the first time.
 - This pair of jobs J_i, J_j forms an inversion
- After swapping i and j we get a schedule with one less inversion
- The new swapped schedule has a maximum lateness no larger than that of O

Swapping two consecutive, inverted jobs



- All jobs other than i and j finish at the same time
- The swap does not increase the lateness of job j
- Lateness of job i is $\ell'_i = f'(i) - d_i = f(j) - d_i < f(j) - d_j = \ell_j$
- It follows that the swap does not increase the max lateness

Correctness

- The initial schedule O has at most $\binom{n}{2}$ inversions (all pairs inverted)
- After at most $\binom{n}{2}$ swaps we get an optimal schedule with no inversions

It follows that:

- *The schedule A produced by the greedy algorithm has optimum lateness L*

Outline

- Greedy algorithms
- Problems solved using greedy
 - Scheduling all intervals (KT – chapter 4.1)
 - Scheduling to minimize lateness (KT – chapter 4.2)
 - **Change making problem (CLRS-problem 16-1 page 446)**
 - The knapsack problem (CLRS page 425-427)

Change-making problem

Represent a given amount of money with the fewest number of coins, when the coins available are quarters (25 cents), dimes (10 cents), nickels (5 cents), and pennies (1 cent).

- let q – number of quarters, d – number of dimes
 k – number of nickels, p – number of pennies

Change-making problem

Greedy algorithm:

Make-change(n)

$S = \phi$

$s = 0$

while $s \neq n$

x is the largest coin such that $s + x \leq n$

if no such coin found

 return “no solution found”

$S = S \cup \{\text{a coin of value } x\}$

$s = s + x$

return S

RT = $O(n)$

Change-making problem

Algorithm 2 (n)

$q = \lfloor n/25 \rfloor$ // number of quarters

$n_q = n \bmod 25$

$d = \lfloor n_q/10 \rfloor$ // number of dimes

$n_d = n_q \bmod 10$

$k = \lfloor n_d/5 \rfloor$ // number of nickels

$n_k = n_d \bmod 5$

$p = n_k$ // number of pennies

$RT = \Theta(1)$

Change-making problem

- example: 89 cents = 3Q + 1D + 4P
- the algorithm produces an optimal solution: there is an optimal solution that makes the greedy choice
- not all coin systems can be solved using the greedy algorithm
 - coin system: 25, 10, 6, 1
 - let $n = 12$
 - greedy: $10 + 1 + 1 \rightarrow 3$ coins
 - optimal: $6 + 6 \rightarrow 2$ coins

Change-making problem

Greedy algorithm:

- if $n = 0$, then the optimal solution has no coins
- if $n > 0$, take the largest coin with value $\leq n$

Let c be this coin. Then use one coin c and recursively solve for $(n - c)$ cents

Greedy Choice Property

Greedy choice property: some optimal solution to the change-making problem for n cents includes a coin with value c , where c is the coin with the largest value $\leq n$.

Proof:

let O be an optimal solution

- if O contains a coin c , then done
- if O does not contain a coin c :
 - if $1 \leq n < 5$, then $c = 1$; all solutions use only P
 - if $5 \leq n < 10$, then $c = 5$

If O does not contain N , then it must use only P ; replace $5P$ by $1N \Rightarrow$ better solution

Greedy Choice Property

- if $10 \leq n < 25$, then $c = 10$

If O does not contain D , then it must use only N and P

Then by replacing a value of 10 ($2N$, $1N+5P$, $10P$) with $1D \Rightarrow$ better solution (smaller number of coins)

- if $n \geq 25$, then $c = 25$

If O does not contain Q , then it must use only D , N , and P ;

Then by replacing a value of 25 with $1Q \Rightarrow$ better solution (smaller number of coins)

$$3D \rightarrow 1Q + 1N$$

$$2D + 1N \rightarrow 1Q$$

$$2D + 5P \rightarrow 1Q$$

.... so on

Since there is always an optimal solution that contains the greedy choice \Rightarrow greedy algorithm produces an optimal solution.

Change-making problem

Problem definition: Suppose that the available coins are in the denominations that are power of c , i.e. the denominations are $c^0, c^1, c^2, \dots, c^k$ for some integers $c > 1$ and $k \geq 1$. Show that the greedy algorithm always yields an optimal solution

Greedy algorithm:

- if $n = 0$, then the optimal solution has no coins
- if $n > 0$, take the largest coin with value $\leq n$

Let c^j be this coin. Then use one coin c^j and recursively solve for $(n - c^j)$ cents

$$\text{RT} = O(k)$$

Greedy Choice Property

Greedy choice property: some optimal solution to the change-making problem for n cents includes a coin with value c^j , where c^j is the coin with the largest value $\leq n$.

Proof:

Let O be an optimal solution

Let a_i be the number of coins of denomination c^i used by O

Note that $a_i < c$, otherwise replace c coins c^i with one coin c^{i+1} and improve the solution

- if O contains a coin c^j , then done
- if O does not contain a coin c^j :
then O contains only coins $c^0, c^1, c^2, \dots, c^{j-1}$

Greedy Choice Property

$$c^j \leq n < c^{j+1}$$

$$\sum_{i=0}^{j-1} a_i c^i = n \geq c^j$$

since O is optimal $\Rightarrow a_i \leq c - 1$ for all $i = 0, 1, 2 \dots j-1$

$$\sum_{i=0}^{j-1} a_i c^i \leq \sum_{i=0}^{j-1} (c-1) c^i = (c-1) \sum_{i=0}^{j-1} c^i = (c-1) \frac{c^j - 1}{c - 1} = c^j - 1$$

geometric series

contradiction \Rightarrow greedy algorithm produces an optimal solution

Outline

- Greedy algorithms
- Problems solved using greedy
 - Scheduling all intervals (KT – chapter 4.1)
 - Scheduling to minimize lateness (KT – chapter 4.2)
 - Change making problem (CLRS-problem 16-1 page 446)
 - **The knapsack problem (CLRS page 425-427)**

The knapsack problem

Given:

- n objects and a knapsack
- $i = 1, \dots, n$ object i has a positive weight w_i and a positive value v_i
- the knapsack can carry a weight $\leq W$

Objective: fill the knapsack such that to maximize the value of the included objects, while respecting the capacity constraints.

Two variations:

- **0-1 knapsack problem**: you can only take the whole object → solved optimally using dynamic programming
- **fractional knapsack problem**: you can take fractions of objects → solved optimally using greedy

Fractional knapsack problem

Greedy algorithm:

- **greedy choice:** choose the item with the largest v_i/w_i value
- this greedy choice produces an optimal solution

Greedy algorithm

Algorithm

sort objects in decreasing order of v_i/w_i

for $i = 1$ to n

$x_i = 0$

load = 0

value = 0

$i = 1$

while load < W and $i \leq n$

if $w_i \leq W - \text{load}$

then take whole object i , $x_i = 1$

else take $x_i = (W - \text{load})/w_i$ of item i

load = load + $x_i w_i$

value = value + $x_i v_i$

$i = i + 1$

- RT = $O(n \log n)$
- example

The greedy algorithm does not work for the 0-1 knapsack problem

Example:

